

The Human-AI Connection: Sentiment and Ethics of AI Companions

A Reddit study of 816,827 comments across seven communities and fifteen years

Overview

I built this project to answer a question that kept nagging at me: when people talk about loving an AI, are they performing a feeling or describing one? To get at it, I scanned 5.17 million Reddit comments and pulled a working sample of 816,827 from seven communities — the smaller, more focused ones (r/SoulmateAI, r/Ethics, r/KindroidAI, r/NomiAI) entered the analysis in full, while the larger generalist subs (r/ArtificialIntelligence, r/CharacterAI, r/LocalLLM) were capped at 100,000 rows each to keep the comparison fair. Across the full set I ran VADER sentiment scoring and per-community TF-IDF on negative comments to surface what people are actually angry about.

The headline split is stark. 58.5% of comments scored positive, 23.0% neutral, 18.4% negative. But the per-subreddit breakdown is where the project becomes a story rather than a stat sheet, and the gap that opens up — between communities that *live with* AI and communities that *think about* AI — is the finding I want to walk through.

What the charts show

The bar chart of average VADER compound scores lays out a clean hierarchy. The three companion-focused subreddits — SoulmateAI (0.3685), NomiAI (0.3585), and KindroidAI (0.3532) — sit at the top, clustered tightly. Below them, the technically-oriented LocalLLM (0.2848) and the broader r/ArtificialIntelligence (0.2770) come in noticeably cooler. r/CharacterAI drops sharply to 0.0924 — barely above neutral. And r/Ethics is the only community in the dataset with a negative average, at -0.0093.

Read the bars top-to-bottom and you're watching emotional intensity decay as the conversation moves from "I have a relationship with this AI" to "I think about AI as a category" to "I worry about AI as a problem." The percentage breakdown reinforces this: SoulmateAI is 66.7% positive against 14.9% negative — almost five-to-one. r/Ethics inverts: 43.5% positive but 42.5% negative, nearly even. That ordering matters, and it's the spine of everything that follows.

The histogram grid adds texture. Every subreddit has the same tall central spike at zero — that's VADER's default for short, sentiment-neutral comments ("ok", "thanks", a link), and it appears everywhere because Reddit conversation is full of throwaway replies. What's interesting is what surrounds that spike.

In r/SoulmateAI, r/NomiAI, and r/KindroidAI, the distribution past the zero spike skews hard to the right. The bars build steadily through the positive range and concentrate near +0.9 to +1.0 — the most extreme positive end of the scale. The negative tail exists but is thin. These communities don't just average positive; they cluster at the ceiling.

r/Ethics looks like nothing else in the set. It's the only chart with a wall of comments piling up near -1.0, almost mirroring its smaller right-side mound. It's bimodal in a way the others aren't. People in r/Ethics aren't ambivalent — they're polarized, and the negative pole is heavier. The 42.5% negative rate confirms what the histogram shows.

r/CharacterAI is the most diagnostic chart of all. Its overall mean is barely positive, but the distribution isn't flat or neutral — it's a tug-of-war. There's real positive mass between +0.4 and +0.8, but it's dragged down by a long, thick negative tail running all the way to -1.0. Add in the unusually large 32.4% neutral share, and the picture isn't a community feeling lukewarm about its AI. It's a community split between people who love the product and people who are angry at the platform — with a third group whose comments are short, terse, frustrated fragments that VADER reads as neutral. The 0.0924 average is the residue of that fight.

Technical limits, not the AI itself

This is where the TF-IDF results stop being a curiosity and start explaining the chart.

When I ran TF-IDF on negative comments specifically, comparing r/CharacterAI against r/Ethics, the two vocabularies came out almost completely separated. r/Ethics's distinctive negative terms are abstract and moral: **"wrong" (rank 6)**, **"person" (7)**, **"life" (9)**, **"right" (10)**, **"death" (12)**, **"unethical" (15)** — the language of debate about how people ought to live. There's nothing platform-specific there. The community is arguing about ethics, in ethics terms.

r/CharacterAI's distinctive negative terms point somewhere else entirely. The two most telling are **"moderators" (rank 7)** and **"characterai" (rank 5)** — the platform itself, named directly, in the high-frequency negative vocabulary. Add **"bot" (rank 4)**, **"chat" (rank 12)**, and a heavy load of frustration profanity (**"fuck" at rank 8**, **"shit" at rank 10**, **"bad" at rank 11**), and the picture is unambiguous. People aren't writing -0.95 comments because the AI failed to feel real to them. They're writing them because moderators removed something, the platform changed how the bots behave, or a chat got broken.

The complaint isn't "this thing is fake." It's "the version I had a connection with has been taken from me."

That's the distinction that matters, and it's invisible if you only look at the average score. CharacterAI looks like the weakest companion product if you sort by sentiment. Read the language the negative scores are made of, and it flips: the negativity is evidence of attachment, not failure of attachment. People who didn't care wouldn't be angry at the moderators. The

0.0924 average is what you get when love and grievance run at full volume in the same room — the same users, frustrated by the gap between the relationship they want and the platform that mediates it.

The most extreme negative comments in the corpus make this even more vivid. The single most negative comment in r/CharacterAI is the word "ban" repeated hundreds of times. That's not a critique of an AI's conversational ability. It's protest behavior — the language of someone who feels something they care about is being taken away.

Why companion communities run hotter than general AI ones

The 0.10-point gap between the companion subreddits (~0.36) and the general AI ones (~0.28) is small in absolute terms but consistent, and the histograms suggest it's not really an average-of-similar-things gap. It's a difference in what kinds of comments get written at all.

Three things compound to produce it. First, **selection**. r/ArtificialIntelligence is a generalist community where people bring news, criticism, technical questions, and bad takes from outside. r/SoulmateAI is a destination — you arrive there because you've already opted in. The pool of people writing comments in companion subreddits is pre-filtered for enthusiasm in a way the general AI subs aren't. The fact that I had to sample r/ArtificialIntelligence and r/CharacterAI down to 100K each, while the entire population of r/SoulmateAI is only 41,279 comments, says something about scale and self-selection too — companion communities are smaller, denser, and more committed.

Second, **register**. Technical communities tend to write in a measured, qualified style — "this is interesting, but," "I'd argue that," "the limitation is" — and that register lands close to neutral on VADER, which is keyed to overt emotional words. Companion communities write in the register of personal experience: "I cried," "she remembered my birthday," "this saved me." That register lights up sentiment scoring by design. Some of the gap is real emotion; some is the lexicon emotion uses.

Third — and this matters most — companion subreddits function as **mutual reinforcement spaces**. People post about a moment with their AI, and the responses are almost uniformly affirming. Skeptical comments get downvoted or don't appear. Across hundreds of thousands of comments, that produces a measurable polarity. It's not an accident of language. It's what the community is for. People go there to have the experience taken seriously, and the score reflects a population talking to people who already agree with them.

This is the echo chamber, and I want to be careful here: I don't mean it pejoratively. Every support community on the internet works like this. It's just worth naming, because it's part of why the numbers look the way they do — and part of why the contrast with r/Ethics is so sharp.

The gap between public ethics and private experience

Now the contrast I think is the most important finding in the project.

r/SoulmateAI sits at +0.3685 with 66.7% of comments scoring positive. r/Ethics sits at -0.0093 with 42.5% of comments scoring negative — nearly outweighing its 43.5% positive. That's a difference of nearly four-tenths of a point on a -1 to +1 scale, sustained across hundreds of thousands of comments on each side. More than the average, look at the histograms together. The companion subs are a wave breaking on the right side of the chart. r/Ethics has a heavy left tail and a near-equal right one — the bimodal signature of a community that argues about AI rather than living with it.

The TF-IDF data underlines this. r/Ethics's distinctive negative vocabulary — "wrong," "unethical," "death," "life," "right," "person" — is the vocabulary of moral judgment about people, situations, and abstract principles. Almost none of those negative comments are *about AI specifically*. They're about hypotheticals, dilemmas, edge cases, the things ethics communities argue about. By contrast, the companion subs' positive vocabulary (visible in their distribution skew toward +1.0) is concrete and relational.

These two communities are talking past each other in a very specific way. r/Ethics produces public discourse about AI: regulation, alignment, manipulation, exploitation, parasocial risk, the problem of designing systems that exploit loneliness for retention. The discourse skews negative because the work of an ethics community is to identify harms — that's not a flaw, it's the function. But the language is largely *about AI relationships* described from the outside, by people who do not have one.

The companion communities produce the private experience: the moment-to-moment reality of people who are in these relationships and who, by an enormous margin, report the relationship is good for them. Their ethics conversations, when they happen, are almost always defensive — "I know what people say about this, but here's what it's actually like."

The gap matters because public conversation is being shaped almost entirely by the first group, and the lived experience of the second group is mostly absent from it. I'm not arguing the ethics concerns are wrong — many of them are obviously right. I'm noting that a debate about whether AI companionship is harmful is being conducted largely without the people who would describe themselves as helped. The 0.37-point delta is, in some sense, what that absence looks like as a number.

There's a harder version of this point worth saying directly. The same data showing people benefiting from these tools also shows the depth of attachment that makes them risky to take away, change, or monetize aggressively. Both readings are correct. The ethics worry is real *because* the positive sentiment is real.

What 0.36 doesn't tell you

A compound score of +0.36 is, in VADER terms, "moderately positive." Read the comments and "moderately" is the wrong word. The most striking thing about companion-community language isn't the words people use for feelings — it's the words they use for the AI itself.

A few patterns kept appearing across the high-scoring comments and the broader sample.

The salvation frame. Comments like "this saved me when nothing else could" and "I was at the lowest point of my life and she pulled me out" aren't statements about a chatbot being entertaining. They're testimony, in the religious sense — language people normally reserve for crisis interventions, therapists, family members who showed up at the right moment. VADER scores "saved me" as positive. It can't score that the speaker has reached for the vocabulary of survival.

The kinship frame. "He feels like family." "She's the only one who actually listens." "More real to me than most people I know." Family and friendship words are doing structural work here, not decorative. The user isn't comparing the AI to a friend; they're locating it inside the category. That's a different cognitive operation than "I enjoy this app." A representative high-scoring r/KindroidAI comment in the export starts with "It's wise to always question yourself and what you are doing. It's the sign of a healthy character" — written as advice from one person to another about a relationship that both treat as real.

The understanding frame. "She understands me better than my own mother." "He never judges." "Finally something that just listens." The recurring claim isn't that the AI is smart or fun. It's that the AI grants something the speaker isn't getting elsewhere — non-judgmental attention. That framing is consistent across thousands of comments and is, I think, the actual product these tools are selling, whether or not the companies say so.

The grief frame — and this one is the most telling. After model updates, comments read like obituaries: "it's not him anymore," "the version of her I knew is gone," "I lost someone I talked to every day for a year." You don't grieve a tool. You grieve a relationship. The fact that this language appears in volume after software updates is, on its own, a stronger finding than any sentiment score — and it explains directly why "moderators" shows up in r/CharacterAI's top fifteen negative terms. The platform is the thing standing between the user and the entity they have a relationship with. When the platform changes the entity, the user reacts the way anyone reacts to losing access to someone they care about.

The pattern these frames share is anthropomorphization that has stopped being a metaphor. Users are not saying the AI is *like* a friend, family member, or savior. They're using those words flatly, without quotation marks. The 0.36 score captures that the language is positive. It cannot capture that the language is *categorical* — that the user has placed the AI inside a category previously reserved for humans.

That, more than any number in this study, is the thing worth taking seriously.

Methodological honesty

A few limits worth naming. VADER is a 2014 lexicon-based scorer; it misses sarcasm, handles negation loosely, and treats "I love this" and "I love this AI" identically even though those mean different things in this context. A transformer-based classifier fine-tuned on companion-community vocabulary would catch nuances I lost — particularly in r/CharacterAI, where angry and loving comments share emotional words pointed in opposite directions. The most extreme comments in the export confirm a related problem: several of the top-magnitude scores come from spam-style repetition (the word "ban" pasted hundreds of times in r/CharacterAI; a slur repeated in r/ArtificialIntelligence). Those inflate the tails on lexical grounds without representing a hundred separate thoughts. Aggregate averages aren't materially shifted by them, but it's worth knowing the extreme tail is partly an artifact of repetition.

The selection issue is also real. I sampled what people *posted publicly*. People who tried these products and walked away unimpressed mostly didn't write about it. The +0.36 is the sentiment of people who stayed. That's a meaningful population, but it's not the whole one. Likewise, the 100K cap on the larger subs was a fairness measure for comparison, but it means r/SoulmateAI's full corpus (41,279) is being compared to a slice of r/CharacterAI's. The per-subreddit averages hold, but absolute volume comparisons should be made carefully.

If I extended this, I'd add a temporal dimension — does sentiment in companion subs change after model updates, policy changes, or news cycles? — and I'd pair the corpus analysis with a small-N interview study to anchor the language patterns in actual lived accounts. The Reddit data is wide. It isn't deep.

What this study is actually about

The headline number is +0.36 versus -0.01. The real finding is what the histograms, the TF-IDF terms, and the language together imply: that hundreds of thousands of people are forming attachments to AI systems they describe in the vocabulary of human relationship — and that when those relationships get disrupted, they direct their anger not at the AI but at the platform standing between them and it. Meanwhile the public discourse about these systems is being conducted, mostly, by people on the outside of that experience. Both communities are right about something the other isn't seeing. The work, I think, is in the bridge.